



ADVANCED OPERATING SYSTEMS AND NETWORKS

Computer Science Engineering

Universidad Complutense de Madrid

2.4. Socket Programming

PROFESSORS:

Eduardo Huedo Cuesta
Rubén Santiago Montero


OTHER AUTHORS:

Rafael Moreno Vozmediano
Juan Carlos Fabero Jiménez


Introduction

Sockets

- A **socket** is an endpoint for communication
- Sockets allow the bidirectional data exchange between a client and a server
- Each server or client application is (normally) identified by a **port** number



- **Socket APIs:** Uniform interface between the user process and the network protocol stacks in the kernel
 - **BSD Sockets API** or Berkeley Sockets
 - WinSock API equivalent to BSD
 - Bindings available for all languages



Sockets Types

- **SOCK_STREAM**

- Sequenced, reliable, two-way, connection-based byte streams
- An out-of-band data transmission mechanism may be supported
- Similarly to a pipe, a connection must be established before any data may be sent or received and the SIGPIPE signal is raised if a process sends or receives on a broken stream
- Message boundaries in incoming datagrams are not preserved, so the beginning and end of the message must be marked (e.g. <HTML></HTML>, {"msg": {...}})

- **SOCK_DGRAM**

- Datagrams (connectionless, unreliable messages of a fixed maximum length)

- **SOCK_RAW**

- Raw network protocol access (e.g. to implement protocol modules in user space)

Socket Creation

- Create an endpoint for communication:

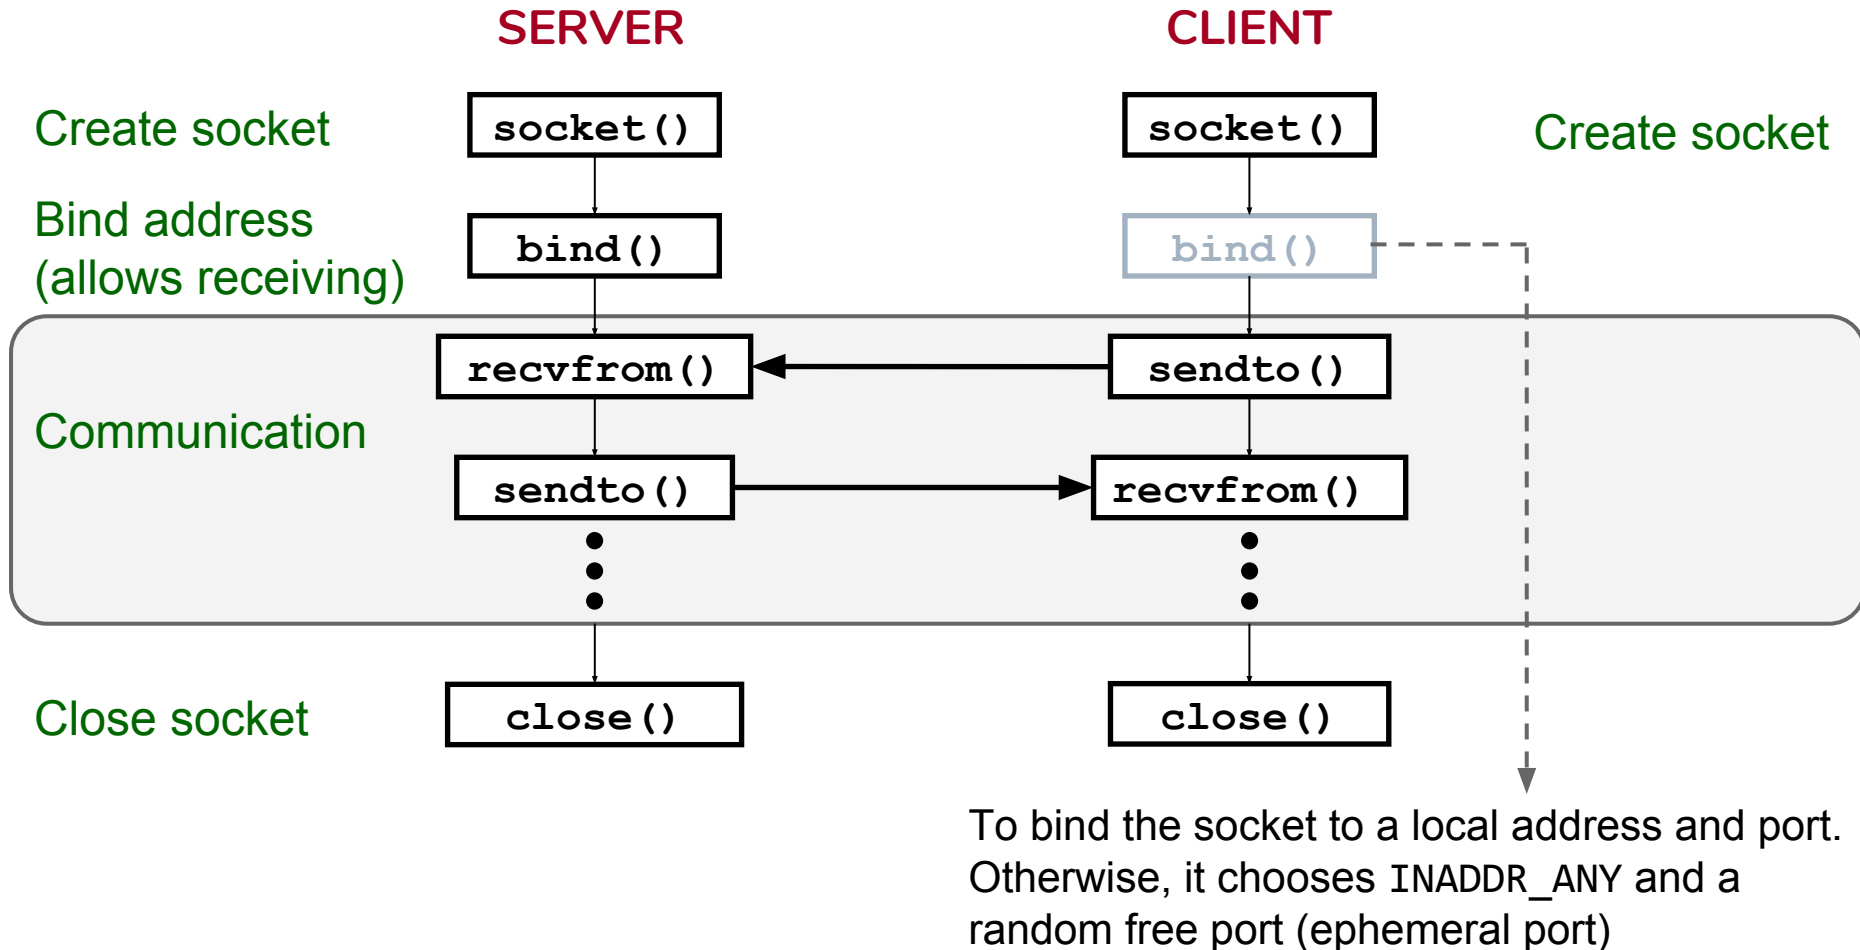
```
int socket(int domain, int type, int protocol);
```

<code><sys/types.h></code> <code><sys/socket.h></code>
POSIX+BSD

- domain is a protocol family, sharing some address scheme
 - AF_UNIX: Local communication between processes in the same system
 - AF_INET, AF_INET6: Internet protocols, over IPv4 or IPv6
 - Others: AF_IPX, AF_X25, AF_APPLETALK, AF_PACKET...
- type is a socket type, specifying the communication semantics
 - SOCK_STREAM, SOCK_DGRAM or SOCK_RAW
- protocol is the particular protocol of the family to be used to implement the socket type
 - Normally, only a single protocol exists to support a particular socket type within a given protocol family, in which case protocol can be specified as 0
 - Some socket types may not be implemented by all protocol families

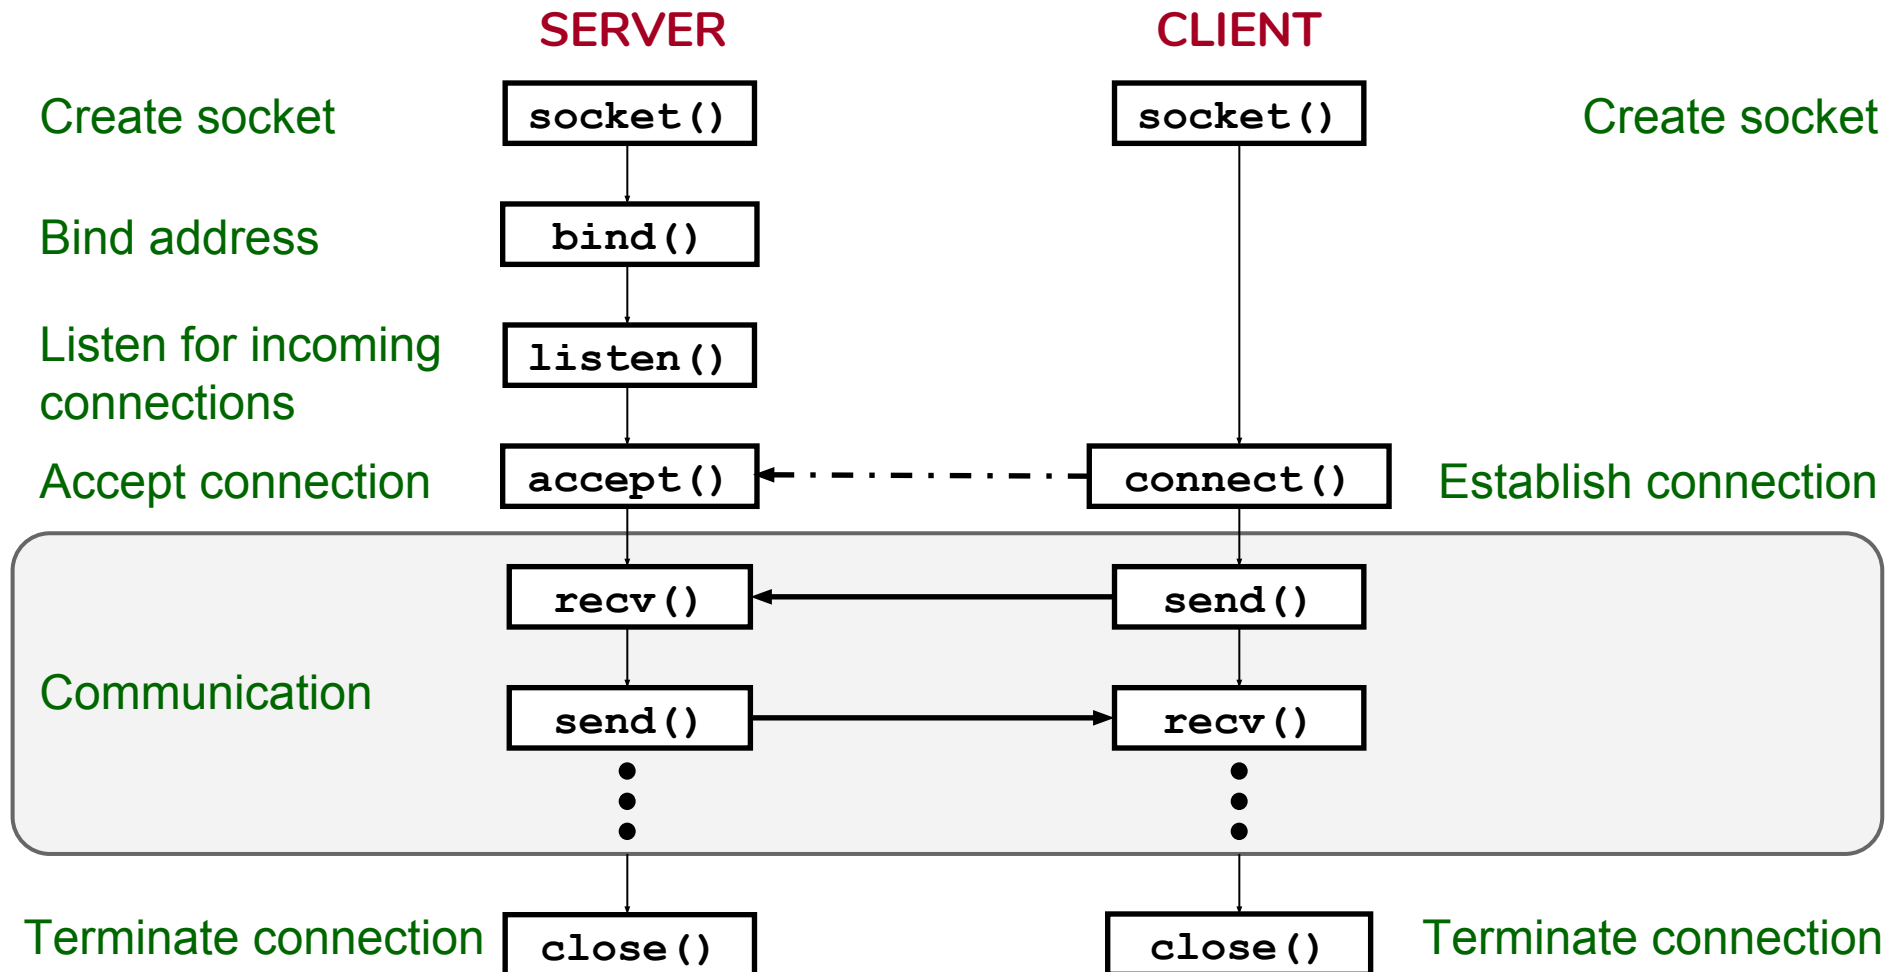
UDP Sockets: Communication Pattern

- Sockets of type `SOCK_DGRAM`



TCP Sockets: Communication Pattern

- Sockets of type SOCK_STREAM



Sockets Creation: IPv4 Sockets

```
tcp_socket = socket(AF_INET, SOCK_STREAM, 0);  
udp_socket = socket(AF_INET, SOCK_DGRAM, 0);  
raw_socket = socket(AF_INET, SOCK_RAW, protocol);
```

- The protocol can be:
 - IPPROTO_TCP for SOCK_STREAM ⇒ always use 0
 - IPPROTO_UDP for SOCK_DGRAM ⇒ always use 0
 - A valid IANA IP protocol number (RFC 1700) for SOCK_RAW
- **Socket address:** Defined by a 32-bit IP address (network-level address) and a 16-bit TCP/UDP port (transport-level address)

```
struct sockaddr_in {  
    sa_family_t    sin_family; // address family: AF_INET  
    in_port_t      sin_port;   // port number in network byte order  
    struct in_addr sin_addr;    // IP address  
};  
  
struct in_addr {  
    uint32_t        s_addr; // 32-bit IP address in network byte order  
};
```

Sockets Creation: IPv4 Sockets

Ports (`in_port_t`)

- Privileged (usually, also well-known) port numbers, below 1024, can only be used by privileged processes
- Associated to transport layer protocols: TCP and UDP
- On RAW sockets it is set to the Protocol field in the IP header

Addresses (`struct in_addr`)

- Address associated to the interface
- To use broadcast addresses, option `SO_BROADCAST` must be specified in the socket
- Constants:
 - `INADDR_ANY` (`0.0.0.0`)
 - `INADDR_BROADCAST` (`255.255.255.255`)
 - For historical reasons, it has the same effect on `bind()` as `INADDR_ANY`
 - `INADDR_LOOPBACK` (`127.0.0.1`)

Management of Network Addresses and Names

- Name-to-address translation in protocol-independent manner:

```
int getaddrinfo(const char *node,  
               const char *service,  
               const struct addrinfo *hints,  
               struct addrinfo **res);
```

<code><sys/types.h></code> <code><sys/socket.h></code> <code><netdb.h></code>
POSIX

- node specifies the host, and it could be:
 - A host name, which is resolved using `gethostbyname(3)`
 - An IPv4 address in dotted-decimal notation, e.g. “192.168.0.1”
 - An IPv6 address in hexadecimal abbreviated notation, e.g. “fe80::1:2”
 - NULL, which refers to the local host
- service specifies the port, and it could be:
 - A service name from `/etc/services` (see `services(5)`), e.g. “http”
 - A decimal integer, e.g. “80”
 - NULL, to not set the port
- hints sets some search criteria
- res is used to return the list of socket addresses
 - The host have multiple network interfaces or supports several protocols (e.g. IPv4 and IPv6)
 - The service supports several protocols (e.g., telnet tcp/23 and udp/23)

Management of Network Addresses and Names

```
struct addrinfo {  
    int             ai_flags;          // Filtering options (hints)  
    int             ai_family;  
    int             ai_socktype;  
    int             ai_protocol;  
    socklen_t       ai_addrlen;       // Result (res)  
    struct sockaddr *ai_addr;  
    char            *ai_canonname;  
    struct addrinfo *ai_next;  
};
```

- Filtering options (hints argument, other fields must be set to 0 or NULL):
 - ai_family: AF_INET for IPv4, AF_INET6 for IPv6, or AF_UNSPEC for both
 - ai_socktype and ai_protocol: Type and protocol as in socket(2)
 - ai_flags: Additional options, e.g. AI_PASSIVE to return 0.0.0.0 or :: with node == NULL (127.0.0.1 or ::1 otherwise)
- Result (res argument):
 - ai_addr and ai_addrlen: Pointer to a socket address and its length in bytes
 - ai_canonname: Official name of the host if AI_CANONNAME set in ai_flags
 - ai_next: Pointer to next result (linked list)

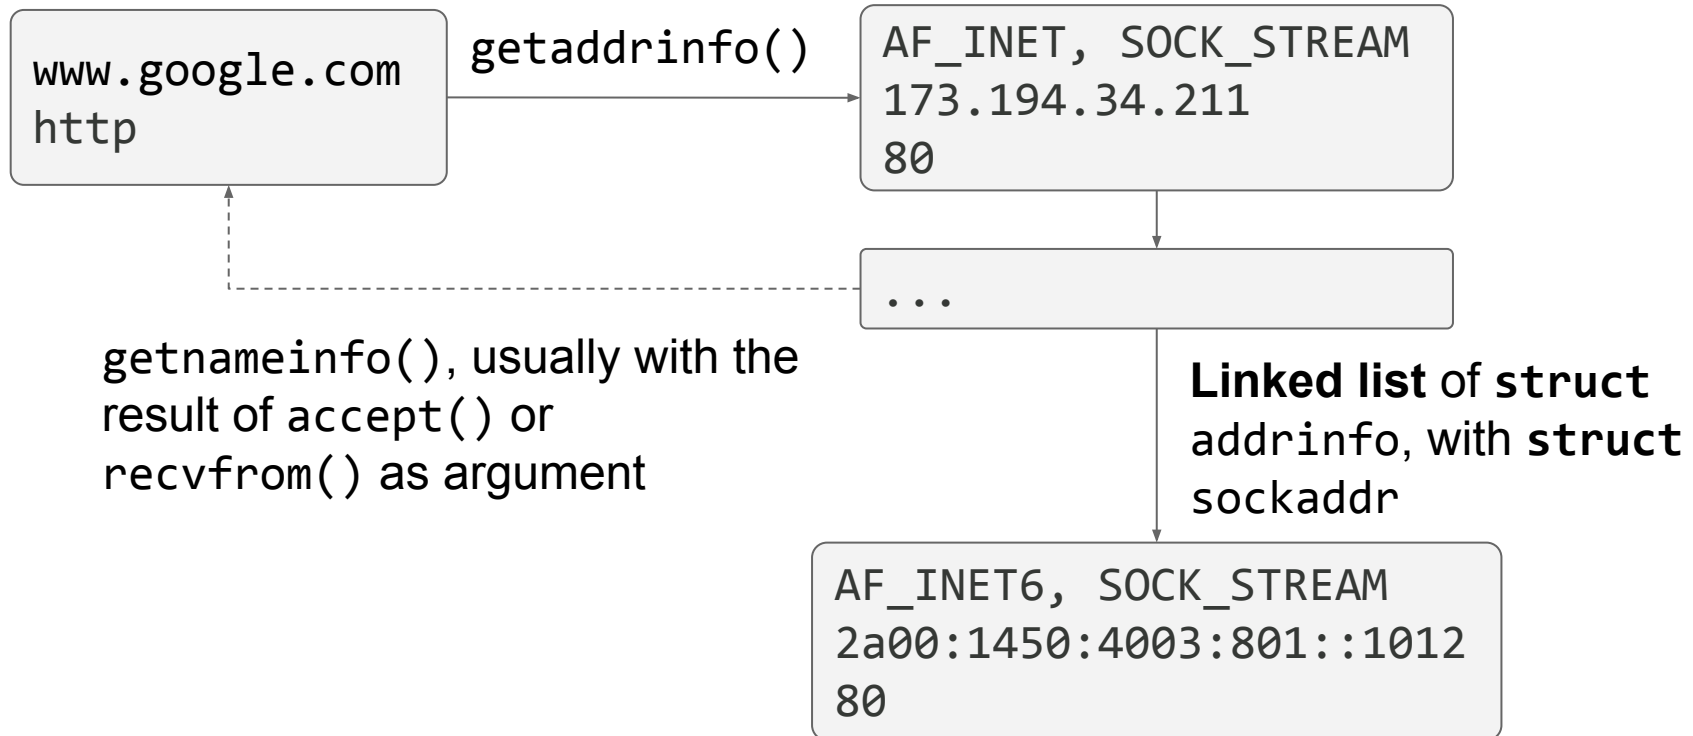
Management of Network Addresses and Names

- Address-to-name translation in protocol-independent manner:

```
int getnameinfo(const struct sockaddr *sa,  
               socklen_t salen, char *host,  
               socklen_t hostlen, char *serv,  
               socklen_t servlen, int flags);
```

```
<sys/types.h>  
<sys/socket.h>  
<netdb.h>
```

POSIX



Management of Network Addresses and Names

- Convert network addresses from binary to text form and from text to binary form:

<arpa/inet.h>
POSIX

```
const char *inet_ntop(int af, const void *src,  
                      char *dst, socklen_t size);  
int inet_pton(int af, const char *src, void *dst);
```

- af is the address family
 - AF_INET: dotted decimal representation (“d.d.d.d”)
 - AF_INET6: abbreviated hexadecimal representation
- Arguments of type **void *** contain the address structure in binary form
 - AF_INET: **struct in_addr**
 - AF_INET6: **struct in6_addr**
- Arguments of type **char *** contain the address representation in text form

Connection Management

- Manage connections:

<code><sys/socket.h></code>


POSIX

```
int bind(int sd, const struct sockaddr *addr,
         socklen_t len);
int connect(int sd, const struct sockaddr *addr, socklen_t len);
int listen(int sd, int backlog);
int accept(int sd, struct sockaddr *addr, socklen_t *len);
```

- sd is a socket descriptor created with a call to `socket()`
- addr is an address of generic type to accommodate each family address
 - In `bind()`, it is the local socket address to be assigned to the socket
 - In `connect()`, it is the address of the socket to connect to
 - In `accept()`, it stores the address of the peer socket after the call returns
- len is the address length, which depends on the address type (use `sizeof()`)
- backlog is the maximum queue length of connections waiting to be accepted
- `listen()` and `accept()` are only used for `SOCK_STREAM`
 - `listen()` marks the socket as a passive socket, that is, as a socket that will be used to accept incoming connection requests using `accept()`
 - `accept()` returns a new socket for the established connection
 - If no pending connections are present on the queue and the socket is not marked as nonblocking, it blocks (`select()` can be used)

Connection Management

```
getaddrinfo(node, service, &hints, &res);
```



```
sock = socket(res->ai_family, res->ai_socktype, res->ai_protocol);
```

```
// Server
```

```
bind(sock, (struct sockaddr *) res->ai_addr, res->ai_addrlen);
```

```
// Client
```

```
connect(sock, (struct sockaddr *) res->ai_addr, res->ai_addrlen);
```

Sending and Receiving Data

- Send and receive data:

<code><sys/socket.h></code>

POSIX

```
ssize_t send(int sock, const void *buf,  
             size_t length, int flags);  
ssize_t recv(int sock, void *buf, size_t length, int flags);
```

- Typically used for SOCK_STREAM
- send() will send length bytes from buf
 - For SOCK_DGRAM, connect() must be used first to set the peer address
- recv() will receive up to length bytes in buf
 - For SOCK_DGRAM, the message should be read in a single recv() operation (buffer size) in order to not losing data
- Both calls can block:
 - Use select() or fcntl(sd, F_SETFL, O_NONBLOCK)
- Errors:
 - EMSGSIZE: Message is too big (the message is not transmitted)

Sending and Receiving Data

- Send and receive data:

```
ssize_t sendto(int socket, const void *buf,  
               size_t length, int flags,  
               const struct sockaddr *dest_addr, socklen_t addrlen);  
ssize_t recvfrom(int socket, void *buf,  
                 size_t length, int flags,  
                 struct sockaddr *src_addr, socklen_t *addrlen);
```

<sys/socket.h>

POSIX

- Typically used for SOCK_DGRAM, to specify or obtain the peer address
- Data is sent in network order, i.e. big-endian, so it may be necessary to convert it to the processor architecture
- Convert values between host and network byte order:

```
uint32_t htonl(uint32_t hostlong)  
uint16_t htons(uint16_t hostshort);  
uint32_t ntohl(uint32_t netlong);  
uint16_t ntohs(uint16_t netshort);
```

<arpa/inet.h>

POSIX

IP Sockets: Options

- Get and set options on sockets:

<code><sys/socket.h></code>

POSIX

```
int getsockopt(int socket, int level,  
               int optname, void *optval, socklen_t *optlen);  
int setsockopt(int socket, int level,  
               int optname, const void *optval, socklen_t optlen);
```

- level specifies where the option is applied:
 - Sockets API: SOL_SOCKET
 - Protocol: IPPROTO_IP, IPPROTO_IPV6, IPPROTO_TCP or IPPROTO_UDP
- Each option (optname) has a specific value (void *optval) and length (optlen). Examples:
 - Socket API: SO_KEEPALIVE, SO_BROADCAST, SO_REUSEADDR
 - IP protocol:
 - Multicast groups: IP_ADD_MEMBERSHIP, IP_DROP_MEMBERSHIP...
 - MTU: IP_MTU, IP_MTU_DISCOVER (Path MTU Discovery)
 - IP datagram fields: IP_OPTIONS, IP_TTL, IP_TOS
- Options can also be changed with sysctl or /proc filesystem (e.g. ip_default_ttl...)

TCP Sockets

- Options `SO_SNDBUF` and `SO_RCVBUF` (Socket API level)
 - Size of send and receive buffers
 - Bigger sizes allow for more efficient management of the TCP window (timestamps, window scaling for flow control...)
 - Available via `sysctl` or `/proc` (`net.ipv4.tcp_rmem` and `net.ipv4.tcp_wmem`)
- Other options (TCP protocol level)
 - `TCP_NODELAY` disables the Nagle algorithm
 - `TCP_QUICKACK` disables delayed ACKs
 - Available via `sysctl` or `/proc`
- Urgent messages (URG flag in TCP header)
 - Add flag `MSG_OOB` (out-of-band) in `send()`
 - The process or process group set as the socket's owner will receive `SIGURG`:
`fcntl(socket, F_SETOWN, pid);`

IPv6 Sockets

```
tcp_socket = socket(AF_INET6, SOCK_STREAM, 0);
udp_socket = socket(AF_INET6, SOCK_DGRAM, 0);
raw_socket = socket(AF_INET6, SOCK_RAW, protocol);
```

- As with AF_INET, SOCK_STREAM is based on TCP and SOCK_DGRAM on UDP
- IPv6 is almost fully compatible with IPv4 implementation
- Socket address:

```
struct sockaddr_in6 {
    sa_family_t      sin6_family;    // Address family: AF_INET6
    in_port_t        sin6_port;      // Port number
    uint32_t          sin6_flowinfo;  // Flow ID
    struct in6_addr   sin6_addr;      // IPv6 address
    uint32_t          sin6_scope_id;  // only if link-local
};

struct in6_addr {
    unsigned char     s6_addr[16];    // 128-bit IPv6 address
};
```

IPv6 Sockets

- The address (`struct in6_addr`) can be instantiated using the following variables:
 - `in6addr_any` (similar to `INADDR_ANY`)
 - `in6addr_loopback` (similar to `INADDR_LOOPBACK`)
- To develop applications compatible with IPv4 and IPv6, dependencies with address format must be avoided
 - `struct sockaddr` is defined just to avoid compiler warnings (same size as `struct sockaddr_in`)
 - `struct sockaddr_storage` is defined to store both `struct sockaddr_in` and `struct sockaddr_in6`
- Example:

```
struct sockaddr_storage addr;  
socklen_t addrlen = sizeof(addr);  
accept(sd, (struct sockaddr *) &addr, &addrlen);
```

- `addr` contains an IPv4 or IPv6 address
- `addrlen` contains the length of the returned structure

IPv4 and IPv6 Servers

- A single socket for IPv6 and IPv4 (dual stack)
 - Disable option `IPV6_V6ONLY` in the socket
 - The default value for this option is defined in `net.ipv6.bindv6only` (disabled by default)
 - Call `bind()` with `in6addr_any` (i.e., `::`)
 - Use of mapped IPv4 addresses (`192.168.0.1` \rightarrow `::FFFF:192.168.0.1`)
 - Not supported in all systems
- Two sockets, one for each stack
 - Enable option `IPV6_V6ONLY` in the IPv6 socket
 - Get valid addresses to create a socket with each stack

Summary: UDP Server Scheme

```
hints.ai_flags      = AI_PASSIVE;    // Return 0.0.0.0 or ::
hints.ai_family     = AF_UNSPEC;     // IPv4 or IPv6
hints.ai_socktype   = SOCK_DGRAM;

rc = getaddrinfo(argv[1], argv[2], &hints, &result);
sd = socket(result->ai_family, result->ai_socktype, 0);
bind(sd, (struct sockaddr *) result->ai_addr, result->ai_addrlen);

while (1) {
    bytes = recvfrom(sd, buf, 80, 0, (struct sockaddr *) &client, &len);
    buf[bytes] = '\0';

    getnameinfo((struct sockaddr *) &client, len, host, NI_MAXHOST,
        serv, NI_MAXSERV, NI_NUMERICHOST|NI_NUMERICSERV);
    printf("Message from %s:%s: %s\n", host, serv, buf);
    sendto(sd, buf, bytes, 0, (struct sockaddr *) &client, len);
}
```

Summary: TCP Server Scheme

```
hints.ai_flags      = AI_PASSIVE;   // Return 0.0.0.0 or ::
hints.ai_family     = AF_UNSPEC;    // IPv4 or IPv6
hints.ai_socktype   = SOCK_STREAM;

rc = getaddrinfo(argv[1], argv[2], &hints, &result);
sd = socket(result->ai_family, result->ai_socktype, 0);
bind(sd, (struct sockaddr *) result->ai_addr, result->ai_addrlen);
listen(sd, 5);

while (1) {
    clisd = accept(sd, (struct sockaddr *) &client, &len);

    getnameinfo((struct sockaddr *) &client, len, host, NI_MAXHOST,
        serv, NI_MAXSERV, NI_NUMERICHOST|NI_NUMERICSERV);
    printf("Connection from %s:%s\n", host, serv);

    while (bytes = recv(clisd, buf, 80, 0)) { // Check message!
        buf[bytes] = '\0';
        printf("\tMessage: %s\n", buf);
        send(clisd, buf, bytes, 0);
    }
}
```

Concurrent Servers

Need

- The server must attend several clients concurrently
- In general, calls are blocking
 - `accept()` waits for the client connection
 - `recv()` and `recvfrom()` wait for data to arrive
 - `send()` and `sendto()` wait if the send buffer is full

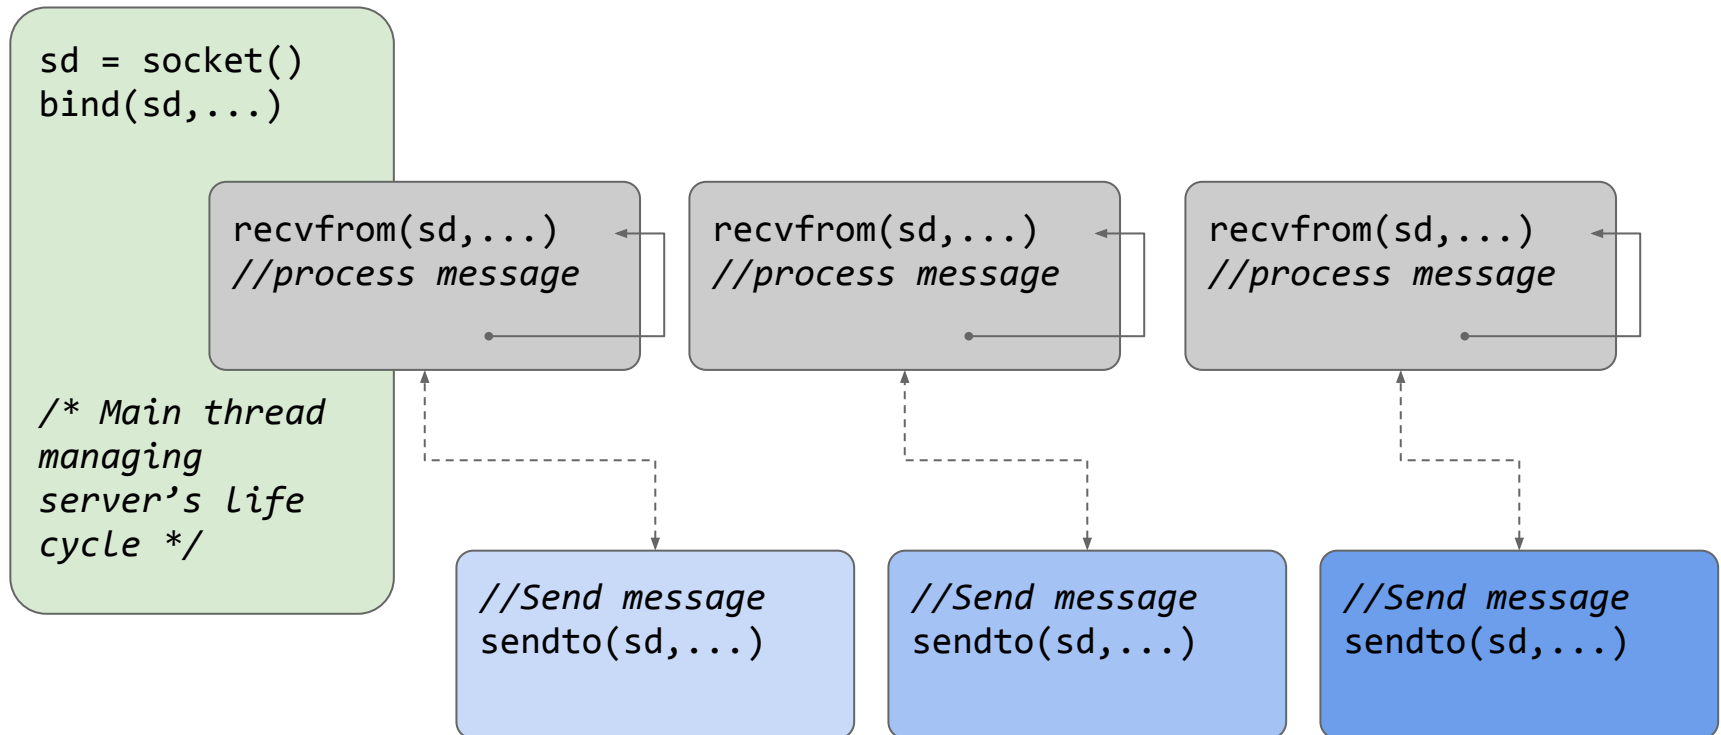
Tools

- Operations are concurrent over socket descriptors
 - Multiple threads can call `accept()` to establish a connection
 - Multiple threads can call `recvfrom()` to receive data
- All threads will block on the call and only one of them will be unblocked when a connection request or data is received
- Threads share address space and file descriptors (sockets)
- Processes inherit file descriptors (sockets)
- Synchronous I/O multiplexing can be also used, but the program logic is much more complex

Concurrent Servers: SOCK_DGRAM

Pre-fork Pattern

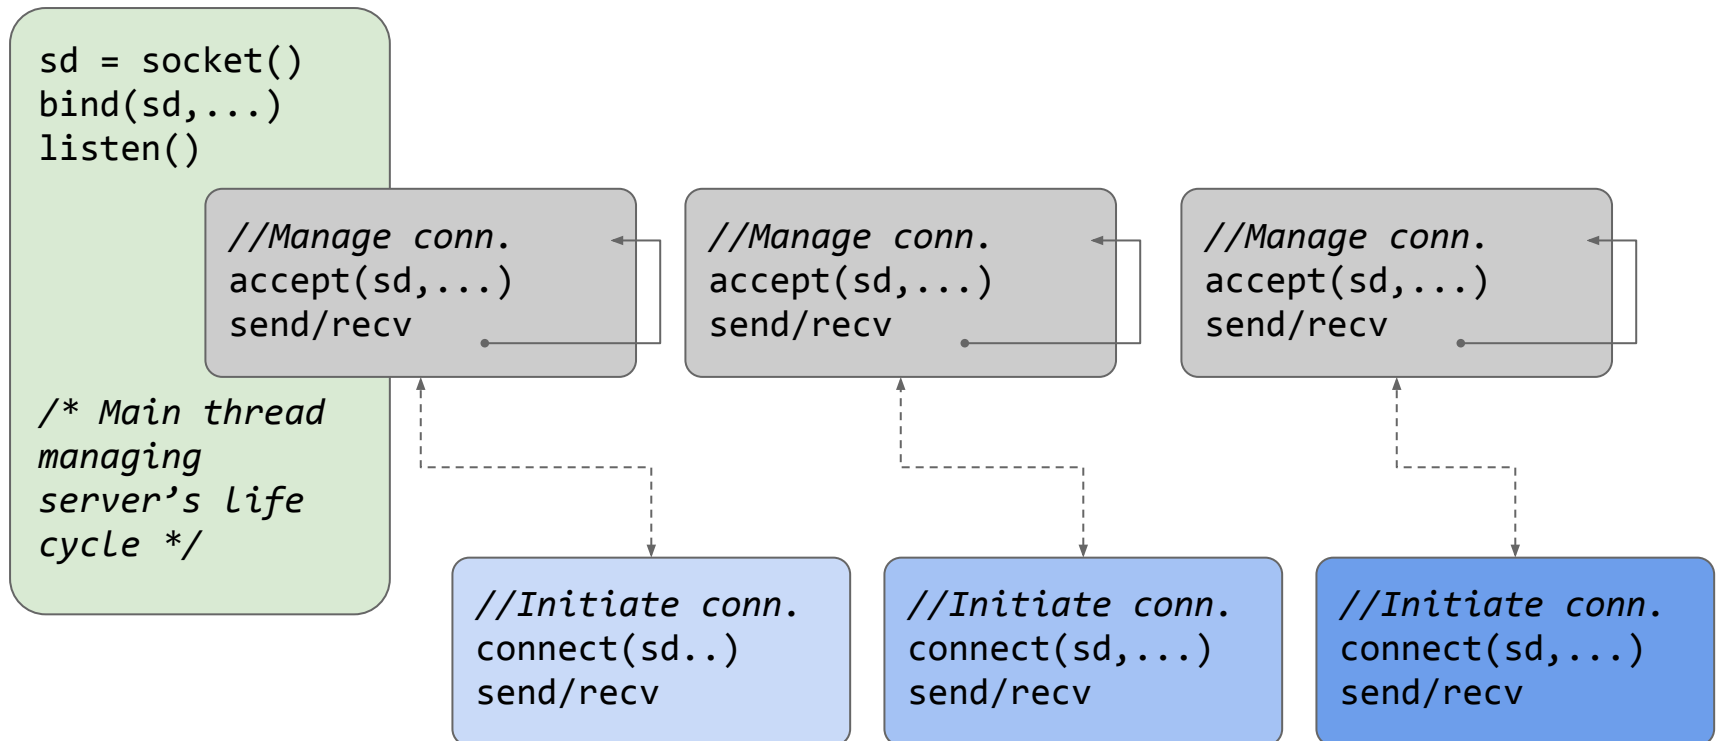
- Concurrent reception of messages
- The server creates a set of processes/threads to receive messages with `recvfrom()` and process them
- Concurrency is at the message level



Concurrent servers: SOCK_STREAM

Pre-fork Pattern

- Concurrent connection management
- The server creates set of processes/threads to manage each connection with `accept()`
- Concurrency is at the connection level



Concurrent Servers: SOCK_STREAM

Accept-and-fork pattern

- Concurrent connection management
- Main process/thread establishes connections with `accept()` and creates a process/thread to manage each one
- Concurrency is at the connection level

